

hw07

November 27, 2018

1 Homework 7: Percentiles, Bootstrap, A/B Testing

1.1 Due Sunday November 25th, 11:59pm

Directly sharing answers is not okay, but discussing problems with the course staff or with other students is encouraged.

You should start early so that you have time to get help if you're stuck.

```
In [167]: #:
import math
import numpy as np
from datascience import *
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
%matplotlib inline

from client.api.notebook import Notebook
ok = Notebook('hw07.ok')
_ = ok.auth(inline=True)
```

```
=====
Assignment: Homework 7
OK, version v1.13.11
=====
```

Successfully logged in as wec149@ucsd.edu

Important: The ok tests don't usually tell you that your answer is correct. More often, they help catch careless mistakes. It's up to you to ensure that your answer is correct. If you're not sure, ask someone (not for the answer, but for some guidance about your approach).

Once you're finished, you must do two things:

1.1.1 a. Turn into OK

Select "Save and Checkpoint" in the File menu and then run the submission cell at the end of the notebook. The result will contain a link that you can use to check that your assignment has been submitted successfully. If you submit more than once before the deadline, we'll only grade your final submission.

1.1.2 b. Turn PDF into Gradescope

Select File > Download As > PDF via LaTeX in the File menu. Turn in this PDF file into the respective assignment at <https://gradescope.com/>. If you submit more than once before the deadline, we will only grade your final submission

1.2 1. Percentiles

The General Definition

Let p be a number between 0 and 100. The p th percentile of a collection is the smallest value in the collection that is *at least as large as* $p\%$ of all the values.

By this definition, any percentile between 0 and 100 can be computed for any collection of values and is always an element of the collection. Suppose there are n elements in the collection. To find the p th percentile:

1. Sort the collection in increasing order.
2. Find $p\%$ of n : $\frac{p}{100} * n$. Call that h . If h is an integer, define $k = h$. Otherwise, let k be the smallest integer greater than h .
3. Take the k th element of the sorted collection.

Question 1. Assign the number of elements in values to the variable n . Define k as above -- your answer should be an integer. Assign the 65th percentile of the array values to `sixty_fifth_percentile`. You must use the variables provided for you when solving this problem. For this problem only, you may *not* use `percentile()`.

Hint: Using `math.ceil()` will round up a number to the next nearest whole number. `math` has already been imported for you.

```
In [168]: #: don't change the values in this array!
          values = make_array(34, 65, 103, 73, 4, 32, 45, 87, 99, 54)
          values.sort() # This line sorts the array
          values
```

```
Out[168]: array([ 4, 32, 34, 45, 54, 65, 73, 87, 99, 103])
```

```
In [169]: n = values.size
          n
```

```
Out[169]: 10
```

```
In [170]: k = math.ceil(65/100 * n)
          k
```

```
Out[170]: 7
```

```
In [171]: sixty_fifth_percentile = values.item(k-1)
          sixty_fifth_percentile
```

```
Out[171]: 73
```

```
In [172]: _ = ok.grade('q1_1')
```

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

**Question 2.** The table `nfl_players` contains one row for each player in the National Football League (NFL). The columns include each player's Number, Name, Position, Age, Height (in inches), Weight (in pounds), College, and Team. Plot a histogram showing each player's weight. Use the bins provided.

Source: <http://espn.com> and <https://www.pro-football-reference.com/players/salary.htm> on 2/1/2018.

```
In [173]: #: read in the table
```


```
nfl_players = Table().read_table('nfl_players.csv')
```

```
nfl_bins = np.arange(140, 370, 5)
```

```
# nfl_players
```

```
In [174]: # plot your histogram here
```

```
nfl_players.hist('Weight', bins=nfl_bins, unit='pounds')
```



**Question 3.** Find the absolute difference between the 95th percentile of the weights of both place kickers (PK) and offensive tackles (OT) in the NFL and assign it to `absolute_difference`. You may use `percentile()`.

```
In [175]: place_kickers = nfl_players.where('Position', are.equal_to('PK')).sort('Weight', desce
# place_percentile = place_kickers.item(math.ceil(95/100 * place_kickers.num_rows))
place_percentile = percentile(95, place_kickers.column('Weight'))
offensive_tackles = nfl_players.where('Position', are.equal_to('OT')).sort('Weight', d
offensive_percentile = percentile(95, offensive_tackles.column('Weight'))
absolute_difference = abs(place_percentile - offensive_percentile)
absolute_difference
```

Out[175]: 116

```
In [176]: _ = ok.grade('q1_3')
```

~~~~~

Running tests

```
-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

Question 4. In an array `quarterback_quartiles` put the values for the first, second, and third quartiles (in that order) of the heights of all quarterbacks (QB) in the NFL. Make sure your values are in the correct order. You may use `percentile()`.

```
In [177]: quarterback_backs = nfl_players.where('Position', are.equal_to('QB')).sort('Weight', desce
first = percentile(25, quarterback_backs.column('Weight'))
second = percentile(50, quarterback_backs.column('Weight'))
third = percentile(75, quarterback_backs.column('Weight'))

quarterback_quartiles = make_array(first, second, third)
quarterback_quartiles
```

Out[177]: array([216, 224, 231])

```
In [178]: _ = ok.grade('q1_4')
```

~~~~~

Running tests

```
-----
Test summary
  Passed: 1
```

```
Failed: 0
[ooooooooook] 100.0% passed
```

**Question 5.** Find the the weights of the heaviest players at each position. Assign the 50th percentile of these weights to `fiftieth_percentile`. You may use `percentile()`.

```
In [179]: # place_kickers = nfl_players.where('Position', are.equal_to('PK')).sort('Weight', des
# offensive_tackles = nfl_players.where('Position', are.equal_to('OT')).sort('Weight',
# quarter_backs = nfl_players.where('Position', are.equal_to('QB')).sort('Weight', des
everyone = nfl_players.group('Position', np.max).sort('Weight', descending=False)

fiftieth_percentile = percentile(50, everyone)
fiftieth_percentile
```

```
Out[179]: 265
```

```
In [180]: _ = ok.grade('q1_5')
```

```
~~~~~
```

Running tests

```

```

```
Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```

**Question 6.** Shaun surveyed his class to find the total number of siblings each of his classmates has. You can see his findings below in the table `siblings`. For instance, 2 people have 0 siblings, 4 have 1 sibling, and so on. If one of his classmates, Jake, has some number of siblings that falls in the 75th percentile of Shaun's data, how many siblings does Jake have? Assign your answer to the value `jake_siblings`. You may use `percentile()`.

```
In [181]: #: load the data
siblings = Table().read_table('siblings.csv')
siblings
```

```
Out[181]: Siblings | Frequency
0 | 2
1 | 4
2 | 5
3 | 4
4 | 3
5 | 1
8 | 1
```

```
In [182]: whole_class = make_array(0, 0, 1, 1, 1, 1, 2, 2, 2, 2, 2, 3, 3, 3, 3, 4, 4, 4, 5, 8)
 jake_siblings = percentile(75, whole_class)
 jake_siblings
```

```
Out[182]: 3
```

```
In [183]: _ = ok.grade('q1_6')
```

```
~~~~~
```

```
Running tests
```

```
-----
```

```
Test summary
```

```
    Passed: 1
```

```
    Failed: 0
```

```
[oooooooooooo] 100.0% passed
```

## 1.3 2. Thanksgiving Dinner

Suppose you're hosting a Thanksgiving dinner and are responsible for cooking a delicious turkey. When you get home from the grocery store, you notice that the turkey you purchased is a bit on the smaller side. You wonder whether turkeys have always been this small. You decide to investigate.

Ideally, you would want to figure out the exact mean weight of all turkeys that were sold for Thanksgiving this year. However, around [46 million turkeys](#) across the world are consumed each Thanksgiving. Thus, it's simply not feasible to obtain the mean weight of *all* turkeys (i.e. the mean weight of the population).

**Question 1.** Complete the statement below by filling in the blanks.

Therefore, you want to collect a sample of turkeys to obtain a population statistic to estimate the unknown (weight) parameter.

You head back to the grocery store and record the weights of all whole turkeys that are in stock. While you're at it, you also decide to record the weights of all whole chickens that are in stock. When you arrive back home, you decide to use this data as your sample.

**Question 2.** Your data is recorded in a CSV file called `stock.csv`. Read this file into a table named `stock`.

```
In [184]: stock = Table.read_table('stock.csv')
         stock
```

```
Out[184]: Fowl      | Weight (lb)
          Turkey   | 7.07681
          Turkey   | 14.1794
          Chicken  | 3.74413
          Turkey   | 18.915
          Chicken  | 3.02597
          Chicken  | 2.32612
          Chicken  | 2.25437
```

```
Turkey | 12.9614
Chicken | 2.92249
Chicken | 2.1849
... (135 rows omitted)
```

```
In [185]: _ = ok.grade('q2_2')
```

```
~~~~~
Running tests
```

```

Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```

**Question 3.** For now, you only care about the weights of the turkeys. Create a new table with the rows of stock where the value of Fowl is "Turkey". Assign this new table to turkeys.

```
In [186]: turkeys = stock.where('Fowl', are.equal_to('Turkey'))
 turkeys
```

```
Out[186]: Fowl | Weight (lb)
 Turkey | 7.07681
 Turkey | 14.1794
 Turkey | 18.915
 Turkey | 12.9614
 Turkey | 13.7908
 Turkey | 16.2523
 Turkey | 22.8031
 Turkey | 9.31993
 Turkey | 21.1311
 Turkey | 20.9763
 ... (52 rows omitted)
```

```
In [187]: _ = ok.grade('q2_3')
```

```
~~~~~
Running tests
```

```
-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

**Question 4.** Calculate the mean weight of turkeys and assign it to `turkey_mean`.

```
In [188]: turkey_mean = np.mean(turkeys.column('Weight (lb)'))
         turkey_mean
```

```
Out[188]: 15.309642114028406
```

```
In [189]: _ = ok.grade('q2_4')
```

```
~~~~~
Running tests
```

```

Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```

You're done! Or are you? You have a single point estimate for the true mean turkey weight. However, you don't know how uncertain your estimate is and you don't know how much these estimates could vary. In other words, you don't have a sense of how good your estimate is. You may have gotten a particular statistic for one sample, but a completely different one for another.

This is where the idea of resampling via the [bootstrap](#) comes in. Let's assume that our original sample resembles the population fairly well. We can then resample from our original sample to produce even more estimates, which we can then use to produce an interval estimate for the true mean weight of all turkeys.

**Question 5.** Fill out the following code to produce 5,000 bootstrapped estimates for the *mean* weight of turkeys. Store your 5,000 estimates in the `turkey_means` array.


```
In [190]: turkey_means = make_array()
 for i in range(5000):
 resample = turkeys.sample(with_replacement=True)
 resample_mean = np.mean(resample.column('Weight (lb)'))
 turkey_means = np.append(turkey_means, resample_mean)

 turkey_means
```

```
Out[190]: array([15.76264232, 16.19138488, 14.93124912, ..., 15.82285974,
 14.94584709, 15.07813662])
```

```
In [191]: #: This cell displays a histogram of turkey_means
 Table().with_column('Estimated Mean', turkey_means).hist()
```





```
In [192]: _ = ok.grade('q2_5')
```

```
~~~~~
```

```
Running tests
```

```
-----
```

```
Test summary
```

```
Passed: 2
```

```
Failed: 0
```

```
[ooooooooook] 100.0% passed
```

**Question 6.** Using the array `means`, compute an approximate 95% confidence interval for the true mean weight of turkeys. (Compute the lower and upper ends of the interval, named `lower_bound` and `upper_bound`, respectively.)

*Hint:* Use `percentile()`.

```
In [193]: lower_bound = percentile(2.5, turkey_means)
          lower_bound
```

```
Out[193]: 14.149131430655984
```

```
In [194]: upper_bound = percentile(97.5, turkey_means)
          upper_bound
```

```
Out[194]: 16.509760417974054
```

```
In [195]: #: the confidence interval
          print("Bootstrapped 95% confidence interval for the true mean weight of turkeys: [{:f}]"
```

```
Bootstrapped 95% confidence interval for the true mean weight of turkeys: [14.149131, 16.509760]
```

```
In [196]: _ = ok.grade('q2_6')
```

```
~~~~~
```

```
Running tests
```

```

```

```
Test summary
```

```
 Passed: 1
```

```
 Failed: 0
```

```
[ooooooooook] 100.0% passed
```

**Question 7.** Which of the following would make the histogram narrower? Assign either 1 or 2 to q2\_7. 1. Increasing the number of resamples (repetitions of bootstrap). 2. Starting with a larger original sample size.

```
In [197]: q2_7 = 2
 q2_7
```

```
Out[197]: 2
```

```
In [198]: _ = ok.grade('q2_7')
```

```
~~~~~
```

```
Running tests
```

```
-----
```

```
Test summary
```

```
    Passed: 1
```

```
    Failed: 0
```

```
[ooooooooook] 100.0% passed
```

**Question 8.** Suppose you want to find the weight of the heaviest turkey (maximum weight out of the entire population). Would your bootstrap procedure be effective in estimating the weight of the heaviest turkey? Explain your answer below.

The procedure would not be efficient, since by bootstrapping, the results will always not reach the real value of the heaviest turkey's weight. All the results will have less value of the real biggest weight.

**Question 9.** Suppose you're wondering how heavy the average turkey is compared to the average chicken. Using the same bootstrap procedure, compute an approximate 95% confidence interval for the true mean difference in weight between turkeys and chickens. Store your 5,000 estimates in the `difference_means` array. Use the original stock table for this.

mean difference := mean weight of turkey – mean weight of chicken

```
In [199]: # You may need to add lines for additional code!
chickens = stock.where('Fowl', are.equal_to('Chicken'))


difference_means = make_array()
for i in range(5000):
    resample_turkey = turkeys.sample(with_replacement=True)
    resample_mean_turkey = np.mean(resample_turkey.column('Weight (lb)'))
    resample_chicken = chickens.sample(with_replacement=True)
    resample_mean_chicken = np.mean(resample_chicken.column('Weight (lb)'))

    difference_means = np.append(difference_means, resample_mean_turkey - resample_mean_chicken)

difference_means

Out[199]: array([10.33003024, 11.48559581, 12.63271859, ..., 11.75118935,
                11.78688656, 12.64114872])

In [200]: #: This cell displays a histogram of difference_means
Table().with_column('Estimated Mean Difference', difference_means).hist()
```



```
In [201]: _ = ok.grade('q2_9')
~~~~~

Running tests

Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```

**Question 10.** Compute the 95% confidence interval for the mean difference in weights of turkeys and chickens. Assign the left and right endpoints to `left_endpoint` and `right_endpoint` respectively.

```
In [202]: left_endpoint = percentile(2.5, difference_means)
 left_endpoint

Out[202]: 10.997345540765405

In [203]: right_endpoint = percentile(97.5, difference_means)
 right_endpoint

Out[203]: 13.398248730349017

In [204]: #: the confidence interval
 print("Bootstrapped 95% confidence interval for the mean difference in weights of turk
Bootstrapped 95% confidence interval for the mean difference in weights of turkeys and chickens:
```

```
In [205]: _ = ok.grade('q2_10')
~~~~~

Running tests

-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

**Question 11:** Based on your histogram and confidence interval, would you say with high probability that the mean turkey is heavier than the mean chicken? Explain your answer.

Since in all the results of our 5000 resamples result in the turkey being at least 10.9 pounds (confidence level) heavier than the chicken, we can conclude that turkey is definitively heavier than the chicken.

**Question 12.** Would changing the units of weight from pounds to kilograms change your conclusion? Assign a boolean (True if it would and False otherwise) to the name `q2_12`.



**Question 1.** Create a new table called `dsc10` with two columns: `Course` and `Letter Grade`. `Course` will contain the string "DSC 10" for all rows and `Letter Grade` will contain the grades from `dsc10_grades`.

*Hint:* Use list multiplication: `3 * ['element'] = ['element', 'element', 'element']`.

```
In [209]: dsc10 = Table().with_column('Course', len(dsc10_grades)*['DSC 10'], 'Letter Grade', ds
dsc10
```

```
Out[209]: Course | Letter Grade
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
... (92 rows omitted)
```

```
In [210]: _ = ok.grade('q3_1')
```

~~~~~

Running tests

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

Question 2. Similarly, create a new table called `poli5d` with two columns: `Course` and `Letter Grade`. `Course` will contain the string "POLI 5D" for all rows and `Letter Grade` will contain the grades from `poli5d_grades`.

```
In [211]: poli5d = Table().with_column('Course', len(poli5d_grades)*['POLI 5D'], 'Letter Grade', ds
poli5d
```

```
Out[211]: Course | Letter Grade
POLI 5D | A
POLI 5D | A
POLI 5D | A
POLI 5D | A
POLI 5D | A
POLI 5D | A
POLI 5D | A
```

```
POLI 5D | A
POLI 5D | A
POLI 5D | A
... (50 rows omitted)
```

```
In [212]: _ = ok.grade('q3_2')
```

```
~~~~~
```

```
Running tests
```

```
-----
```

```
Test summary
```

```
Passed: 1
```

```
Failed: 0
```

```
[ooooooooook] 100.0% passed
```

Question 3. Create a new table `letter_grades` and append to it both the `poli5d` table and the `dsc10` table.

Hint: `original_tbl.append(new_tbl)` appends `new_tbl` to `old_tbl`, but `old_table` must have all of the columns of `new_tbl`. Try using `letter_grades = Table(['Course', 'Letter Grade'])` to create an empty table with the necessary columns.

Warning: You should be able to run the cell below twice and get the same number of rows each time! If not, you might not be creating a *new* table.

```
In [213]: letter_grades = Table(['Course', 'Letter Grade']).append(dsc10).append(poli5d)
          letter_grades
```

```
Out[213]: Course | Letter Grade
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
DSC 10 | A
... (152 rows omitted)
```

```
In [214]: _ = ok.grade('q3_3')
```

```
~~~~~
```

```
Running tests
```

```
-----
```

```
Test summary
```

```
Passed: 1
Failed: 0
[ooooooooook] 100.0% passed
```

Note that `letter_grades` now consists of all 102 rows for DSC 10 followed by all 60 rows for POLI 5D (or *vice versa*, depending on your particular solution). In other words, the `Course` column is not randomly shuffled; all rows for DSC 10 come before all rows for POLI 5D. In your upcoming permutation test, you will shuffle the `Letter Grade` column but will keep `Course` in the same order.

Question 4. Is it a problem that the `Course` column will never be shuffled in your permutation test? Explain your answer below.

Yes. Because the `'Course'` column is not shuffled along with the `'Letter Grade'` column, each letter grade in the `'Letter Grade'` column will not match its corresponding course in the `'Course'` column. So the results of the experiments would be useless, since the data got mixed up and no meaningful results would come out

There is one last thing we need to take care of before we can start our permutation test. As mentioned earlier, grades are only shown as letter grades, not numerical grades. However, working with numerical grades rather than letter grades will allow us to use the permutation test more effectively. This is because we can take the mean of numerical grades while we cannot with letter grades.

Therefore, we will map the letter grades in our data set to placeholder numerical grades according to the following mapping:

Letter Grade	Placeholder Numerical Grade
A	4.0
B	3.0
C	2.0
D	1.0
F	0.0

Note that these placeholder numerical grades are just approximations of the actual numerical grades received. For example, a student under the "A" category may have actually received an A-, which corresponds to 3.7 grade points, rather than a 4.0. Nevertheless, we will make do with what we have.

Question 5. Define a function `letter_to_numerical` that takes in a letter grade as a string as input, and returns the corresponding placeholder numerical grade according to the mapping above.

```
In [215]: # write your function here
def letter_to_numerical(letter_grade):
    if letter_grade == 'A':
        return 4.0
    elif letter_grade == 'B':
        return 3.0
    elif letter_grade == 'C':
```



```

        return 2.0
    elif letter_grade == 'D':
        return 1.0
    elif letter_grade == 'F':
        return 0.0

```

```
In [216]: _ = ok.grade('q3_5')
```

```
~~~~~
```

Running tests

```
-----
```

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

Question 6. Using your function `letter_to_numerical`, create a new table `numerical_grades` with two columns: `Course` and `Numerical Grade`. `Course` will contain the same values as in the `letter_grades` table and `Numerical Grade` will contain corresponding placeholder numerical grades.

```
In [217]: numberss = letter_grades.apply(letter_to_numerical, 'Letter Grade')
numerical_grades = letter_grades.select('Course').with_column('Numerical Grade', numberss)
numerical_grades
```

```
Out[217]: Course | Numerical Grade
```

```

DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
DSC 10 | 4
... (152 rows omitted)

```

```
In [218]: _ = ok.grade('q3_6')
```

```
~~~~~
```

Running tests

```
-----
```

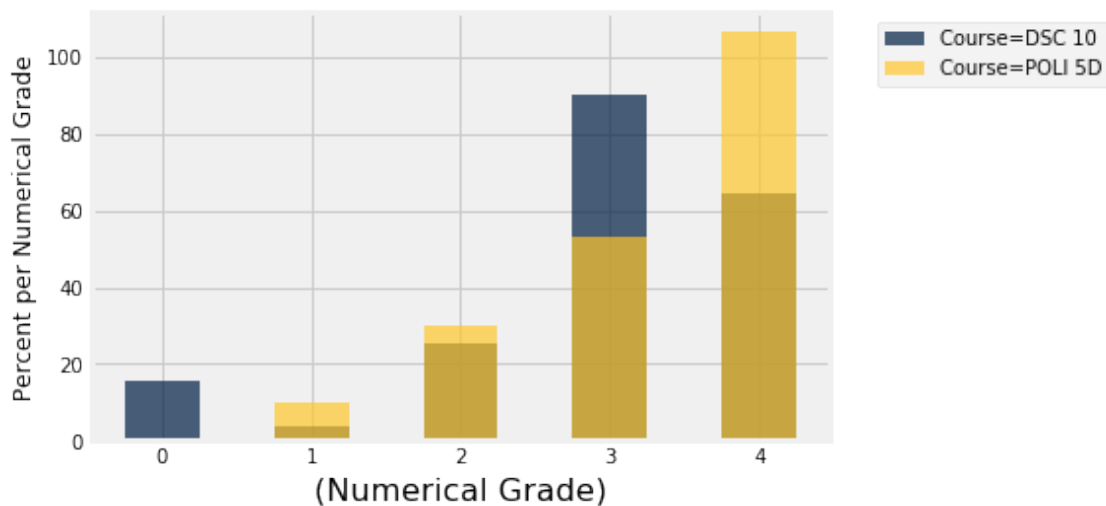
Test summary

Passed: 1

```
Failed: 0
[oooooooooooo] 100.0% passed
```

Question 7. Run the below cell to display a histogram of numerical_grades. Why are there gaps in-between the bars? Explain your answer below.

```
In [219]: #: Just run this cell; do not change it
          numerical_grades.hist('Numerical Grade', unit='Numerical Grade', group='Course', bins=
```



Because in the bins it is specified to hve a 0.5 gap between each bar

Question 8. Using letter_grades, calculate the difference between the mean DSC 10 grade and the mean POLI 5D grade. Assign your answer to observed_difference.

mean difference := mean DSC 10 grade – mean POLI 5D grade

```
In [220]: dsc_mean = np.mean(numerical_grades.where('Course', are.equal_to('DSC 10')).column('Nu
          poli_mean = np.mean(numerical_grades.where('Course', are.equal_to('POLI 5D')).column('
          observed_difference = dsc_mean - poli_mean
          observed_difference
```

```
Out[220]: -0.361764705882353
```

```
In [221]: _ = ok.grade('q3_8')
```

~~~~~

Running tests

-----

Test summary

```
Passed: 1
Failed: 0
[ooooooooook] 100.0% passed
```

**Question 9.** Interpret in words the number you obtained for `observed_difference`. Explain your answer below. Be as specific as possible.

`observed_difference` represents the mean of the grades from dsc 10 minus the mean of the grades from poli 5d. Being it a negative number, it tells us that the mean of the grades of dsc 10 is lower than the mean of the grades from poli 5d. This means that the average grade of dsc 10 is lower than the average grade of poli 5d. Suggesting that dsc 10 has a lower average grade than poli 5d.

**Question 10.** Use a permutation test to calculate 5,000 differences using random permutations of the data. Store your 5,000 differences in the `differences` array.

```
In [222]: differences = make_array()
          for i in range(5000):
              shuffled_grades = numerical_grades.sample(with_replacement=False).column('Numerical Grade')
              shuffled = numerical_grades.with_column('Shuffled Numerical Grade', shuffled_grades)
              group_means = shuffled.group('Course', np.mean).column('Shuffled Numerical Grade mean')

              difference = group_means.item(1) - group_means.item(0)
              differences = np.append(differences, difference)

          differences

Out[222]: array([-0.11470588,  0.04411765,  0.09705882, ..., -0.16764706,
                 0.04411765, -0.16764706])
```

```
In [223]: _ = ok.grade('q3_10')
```

```
~~~~~
Running tests
```

```

Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```


**Question 11.** Which of the follow choices best describes the purpose of the permutation test with regards to A/B testing? Assign either 1, 2, or 3 to `q3_11`. 1. The permutation test generates a null distribution which we can use in testing our hypothesis. 2. The permutation test mitigates noise in our data by generating new permutations of the data. 3. The permutation test is a special case of the bootstrap and allows us to produce interval estimates.

```
In [224]: q3_11 = 2
 q3_11
```

```
Out[224]: 2
```

The following cell plots your observed difference and a histogram of your null distribution produced under your permutation tests.

```
In [225]: #: Just run this cell; do not change it
 Table().with_column('Difference in Mean', differences).hist()
 plt.scatter(observed_difference, 0, color='red', s=40);
```



```
In [226]: _ = ok.grade('q3_11')

~~~~~

Running tests
```

```
-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

**Question 12.** Compute a p-value for the hypothesis. That is, under the null hypothesis, compute the probability that we would have obtained a difference equal to or lower than observed\_difference by chance alone. Assign your answer to p\_val.

```
In [227]: p_val = np.count_nonzero(differences <= observed_difference) / 5000
          p_val
```

```
Out[227]: 0.0178
```

```
In [228]: _ = ok.grade('q3_12')
```

```
~~~~~
Running tests
```

```

Test summary
 Passed: 1
 Failed: 0
[ooooooooook] 100.0% passed
```

**Question 13.** Do you reject or fail to reject the null hypothesis at the 0.05 significance level? What conclusion can you make with regards to DSC 10 and POLI 5D? Explain your answer below. Be as specific as possible.

At the 0.05 significance level, we reject the null hypothesis in favor to the alternative hypothesis. We can conclude that grades received from DSC 10 are typically lower than those from POLI 5D.

**Question 14.** Suppose you are interested in two values, value *A* and value *B*. Suppose your null hypothesis is as follows:

**Null hypothesis:** In the population, value *A* is equal to value *B*.

Suppose your observed difference (value *A* – value *B*) lies in the far left tail of the null distribution.

Which alternative hypothesis will most likely result in a p-value that is NOT significant (a p-value NOT less than 0.05)? Assign either 1, 2, or 3 to q3\_14.

1. **Alternative hypothesis:** There is a difference between value *A* and value *B*.
2. **Alternative hypothesis:** Value *A* is greater than value *B*.
3. **Alternative hypothesis:** Value *B* is greater than value *A*.

```
In [229]: q3_14 = 1
 q3_14
```

Out[229]: 1

In [230]: \_ = ok.grade('q3\_14')

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

To submit:

1. Select Kernel -> Restart & Run All to ensure that you have executed all cells, including the test cells.
2. Read through the notebook to make sure everything is fine and all tests passed.
3. Submit using the cell below.
4. Save PDF and submit to gradescope

In [231]: #: Run all tests at once

import os

\_ = [ok.grade(q[:-3]) for q in os.listdir('tests') if q.startswith('q')]

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
Passed: 1  
Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 2

Failed: 0

[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~



Running tests

-----  
Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~  
Running tests

-----  
Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~  
Running tests

-----  
Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~  
Running tests

-----  
Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~  
Running tests

-----  
Test summary

Passed: 1

Failed: 0

[ooooooooook] 100.0% passed

~~~~~  
Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----

Test summary  
    Passed: 1  
    Failed: 0

[oooooooooooo] 100.0% passed

### 1.5 Before submitting, select "Kernel" -> "Restart & Run All" from the menu!

Then make sure that all of your cells ran without error.

```
In [232]: #: Submit your notebook
 _ = ok.submit()
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.Javascript object>
```

```
Saving notebook... Saved 'hw07.ipynb'.
Submit... 100% complete
Submission successful for user: wec149@ucsd.edu
URL: https://okpy.org/ucsd/dsc10/fa18/hw07/submissions/VP4Jkz
```

### 1.6 Don't forget to submit to both OK and Gradescope!